



Assignment: Modern ML & Neural Networks

Objective

The objective of this assignment is to:

- Understand **classical machine learning models**
- Implement **core algorithms from scratch**
- Compare **manual implementations vs library implementations**
- Analyze **model performance using metrics**
- Practice **clean OOP-based coding** and **experimental reporting**

You will work primarily on a **movie completion prediction dataset**, and partially on a **music genre classification dataset**.

⚠ Academic Integrity

- Plagiarism is strictly prohibited
- You may use libraries **only where allowed**
- All manual implementations must be your own

📁 Datasets

1. Movie Completion Dataset

<https://www.kaggle.com/datasets/cmaigrot/movies-and-series-rating-from-imdb>

2. Music Genre Classification Dataset

<https://www.kaggle.com/datasets/andradaolteanu/gtzan-dataset-music-genre-classification>

🧩 Part 1: Decision Tree & Random Forest

1.1 Decision Tree using Library

As demonstrated in class, use a **Decision Tree classifier from a standard ML library** to predict whether a user completes a movie.

You must:



- Train a Decision Tree model
- Try **at least 3 different hyperparameter combinations**, such as:
 - `max_depth`
 - `min_samples_split`
 - `criterion` (gini / entropy)
- Evaluate each model using:
 - Accuracy
 - F1-score

1.2 Comparison Table (Mandatory)

Create a **tabular comparison** in your report with the following columns:

Model	<code>max_depth</code>	<code>min_samples_split</code>	<code>criterion</code>	Accuracy	F1 Score
-------	------------------------	--------------------------------	------------------------	----------	----------

★ Bonus (Optional): Random Forest from Scratch

- Implement the **Random Forest algorithm manually**
- You **may use Decision Trees from a library**, but:
 - Bootstrapping
 - Feature sampling
 - Aggregation (majority voting)**must be written by you**
- Compare:
 - Decision Tree vs Random Forest
 - Accuracy
 - F1 Score

🧠 Part 2: Logistic Regression

2.1 Logistic Regression from Scratch

- Implement **Logistic Regression manually**
- Use:
 - Gradient Descent
 - Sigmoid activation

- Binary cross-entropy loss
- Follow **OOP principles** (use a class)

Evaluate:

- Accuracy
 - F1 Score
- on the movie dataset.
-

2.2 Logistic Regression using Library

- Train Logistic Regression using a standard ML library
 - Evaluate:
 - Accuracy
 - F1 Score
 - Compare results with your manual implementation
-

★ Bonus (Theory): Ensemble Methods with Logistic Regression

In the report, **theoretically explain**:

Which ensemble method (Bagging, Boosting, Stacking) would be more suitable when using Logistic Regression for this movie dataset, and why?

Support your answer using:

- Bias-variance intuition
 - Dataset characteristics
-

🧠 Part 3: Neural Networks

3.1 Neural Network from Scratch (Movie Dataset)

- Implement a **small Neural Network from scratch**
- Example architecture:
 - 3 layers
 - 5 neurons per layer
- Use:
 - Forward propagation
 - Backpropagation

- Gradient descent
- Follow **OOP design** (class-based implementation)

Train on:

- Movie completion dataset

Report:

- Accuracy
 - F1 Score
-

3.2 Neural Network using Library (Music Genre Dataset)

- Use a deep learning library (TensorFlow / PyTorch)
 - Train a Neural Network on **music genre classification**
 - Evaluate:
 - Accuracy
 - F1 Score
-

3.3 Hyperparameter Tuning (Mandatory)

- Study basic concepts of **hyperparameter tuning**
- Try **at least 5 different hyperparameter combinations**, such as:
 - Learning rate
 - Number of layers
 - Number of neurons
 - Batch size
 - Activation functions

Report:

- Best model configuration
 - Corresponding Accuracy & F1 Score
-



Visualization Requirements (Mandatory)

Wherever model training is involved, include:

- Training loss vs epochs
- Validation loss vs epochs

Plots must be:

- Clearly labeled
 - Properly explained in the report
-



Coding Guidelines (Strict)

- Any **manual implementation** must:
 - Use **Python classes**
 - Follow **OOP principles**
 - Code should be:
 - Clean
 - Modular
 - Well-commented
-



Submission Format

You must submit a zip file containing:

1 Jupyter Notebook (.ipynb)

- All code cells must be:
 - Executed
 - Outputs visible
- Proper section headings
- No broken cells

2 Report (separate PDF)

The report must include:

- Explanation of each model
 - Hyperparameter choices
 - Tables of results
 - Graphs and interpretations
 - Your **methodology and thought process**
-



Evaluation Criteria

Component	Weight
Correct implementation	40%
Experimental comparison	20%
Report clarity & explanation	20%
Code quality & OOP usage	10%
Bonus sections	10%